

# HOW MOLECULES IMPACT CELLS: UNLOCKING CONTRASTIVE PHENOMOLECULAR RETRIEVAL

**Philip Fradkin<sup>1,3,\*</sup>, Puria Azadi<sup>1,2,\*</sup>, Karush Suri<sup>1</sup>,  
Frederik Wenkel<sup>1</sup>, Ali Bashashati<sup>2</sup>, Maciej Sypetkowski<sup>1†</sup>, Dominique Beaini<sup>1,4,†</sup>**

<sup>1</sup> Valence Labs, <sup>2</sup> University of British Columbia,

<sup>3</sup> University of Toronto, Vector Institute, <sup>4</sup> Université de Montréal,

Mila- Quebec AI Institute

dominique@valencelabs.com

## ABSTRACT

Predicting molecular impact on cellular function is a core challenge in therapeutic design. Phenomic experiments, designed to capture cellular morphology, utilize microscopy based techniques and demonstrate a high throughput solution for uncovering molecular impact on the cell. In this work, we learn a joint latent space between molecular structures and microscopy phenomic experiments, aligning paired samples with contrastive learning. Specifically, we study the problem of *Contrastive PhenoMolecular Retrieval*, which consists of zero-shot molecular structure identification conditioned on phenomic experiments. We assess challenges in multi-modal learning of phenomics and molecular modalities such as experimental batch effect, inactive molecule perturbations, and encoding perturbation concentration. We demonstrate improved multi-modal learner retrieval through (1) a uni-modal pre-trained phenomics model, (2) a novel inter sample similarity aware loss, and (3) models conditioned on a representation of molecular concentration. Following this recipe, we propose *MolPhenix*, a molecular phenomics model. MolPhenix leverages a pre-trained phenomics model to demonstrate significant performance gains across perturbation concentrations, molecular scaffolds, and activity thresholds. In particular, we demonstrate an  $8.1\times$  improvement in zero shot molecular retrieval of active molecules over the previous state-of-the-art, reaching 77.33% in top-1% accuracy. These results open the door for machine learning to be applied in virtual phenomics screening, which can significantly benefit drug discovery applications.

## CONTENTS

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Related Work</b>	<b>4</b>
<b>3</b>	<b>Methodology</b>	<b>4</b>
<b>4</b>	<b>Experimental Setup</b>	<b>8</b>
<b>5</b>	<b>Results and Discussion</b>	<b>8</b>

\*These two authors contributed equally and reserve the right to swap their order.

†Equal advising.

5.1	Evaluation on cumulative concentrations . . . . .	8
5.2	Evaluation on held-out concentrations . . . . .	10
5.3	Ablation Studies . . . . .	10
<b>6</b>	<b>Conclusion</b>	<b>11</b>
<b>7</b>	<b>Acknowledgements</b>	<b>11</b>
<b>A</b>	<b>Assumption of the Initial Cell State</b>	<b>17</b>
<b>B</b>	<b>Dataset</b>	<b>17</b>
B.1	Concentration Details . . . . .	17
<b>C</b>	<b>Implementation Details</b>	<b>18</b>
C.1	Hyperparameters . . . . .	18
C.2	Resource Computation . . . . .	18
C.3	S2L Distance function . . . . .	18
<b>D</b>	<b>Additional Results</b>	<b>19</b>
D.1	Predicting molecular activity . . . . .	19
D.2	Zero Shot Biological Validation . . . . .	21
D.3	Addressing Challenges in Contrastive Phenomic Retrieval . . . . .	22
D.4	Ablation Studies . . . . .	22

## 1 INTRODUCTION

Quantifying cellular responses elicited by genetic and molecular perturbations represents a core challenge in medicinal research [Bock et al. \(2022\)](#); [Vincent et al. \(2022\)](#). Out of an approximate  $10^{60}$  druglike molecule designs, a small number are able to alter cellular properties to reverse the course of diseases [Bohacek et al. \(1996\)](#); [Knox et al. \(2023\)](#). In recent years, microscopy-based cell morphology screening techniques, demonstrated potential for quantitative understanding of a molecule’s biological effects. Experimental techniques such as cell-painting are used to capture cellular morphology, which correspond to physical and structural properties of the cell [Boutros et al. \(2015\)](#); [Bray et al. \(2016\)](#). Cells treated with molecular perturbations can change morphology, which is captured by staining and high throughput microscopy techniques. Perturbations with similar cellular impact induce analogous morphological changes, allowing to capture underlying biological effects in phenomic experiments. Identifying such perturbations with similar morphological changes can aid in discovery of novel therapeutic drug candidates [Simm et al. \(2018\)](#); [Kraus et al. \(2017\)](#); [Hofmarcher et al. \(2019\)](#).

Determining molecular impact on the cell can be formulated as a multi-modal learning problem, allowing us to build on a rich family of methods [Radford et al. \(2021\)](#); [Zhai et al. \(2023\)](#); [Srinivasa et al. \(2023\)](#). Similar to text-image models, paired data is collected from phenomic experiments along with molecules used to perturb the cells. Contrastive objectives have been used as an effective approach in aligning paired samples from different modalities [Radford et al. \(2021\)](#); [Lanusse et al. \(2023\)](#). A model that has learned a cross-modal joint latent space must be able to retrieve a molecular perturbant conditioned on the phenomic experiment. We identify this problem as *contrastive phenomolecular retrieval* (see Figure 2). Addressing this problem can allow for identification of molecular impact on cellular function,

however, this comes with its own set of challenges. Fürst et al. (2022); Albelwi (2022); Zhong et al. (2022).

(1) Firstly, multi-modal paired phenomics molecular data suffers from lower overall dataset sizes and is subject to batch effects. Challenges with uniform processing and prohibitive costs associated with acquisition of paired data, leads to an order of magnitude fewer data points compared to text-image datasets Schuhmann et al. (2022); Chandrasekaran et al. (2023). Furthermore, data is subject to random batch effects that capture non-biologically meaningful variation Leek et al. (2010); Sypetkowski et al. (2023). (2) Paired phenomic-molecular data contains inactive perturbations that do not have a biological effect or do not perturb cellular morphology. It is difficult to infer a priori whether a molecule has a cellular effect, leading to the collection of paired molecular structures with unperturbed cells. These data-points are challenging to filter out without an effective phenomic embedding, as morphological effects are rarely discernible. These samples can be interpreted as misannotated, under the assumption of all collected pairs having biologically meaningful interactions. (3) Finally, a complete solution for capturing molecular effects on cells must capture molecular concentration. The same molecule can have drastically different effects along its dose response curve, thus making concentration an essential component for learning molecular impact.

In this work, we explore the problem of contrastive phenomolecular retrieval by addressing the above challenges circumvented in prior works. Our key contributions are as follows:

- We demonstrate significantly higher phenomolecular retrieval rates by utilizing a pre-trained uni-modal phenomic encoder. Thus alleviating the data availability challenge, reducing the impact of batch effects, and identifying molecular activity levels.
- We propose a novel soft-weighted sigmoid locked loss (S2L) that addresses the effects of inactive molecules. This is done by leveraging distances computed in the phenomic embedding space to learn inter-sample similarities.
- We explore *explicit* and *implicit* methods to encode molecular concentration, assessing the model’s ability to perform retrieval in an inter-concentration setting and generalize to unseen concentrations.

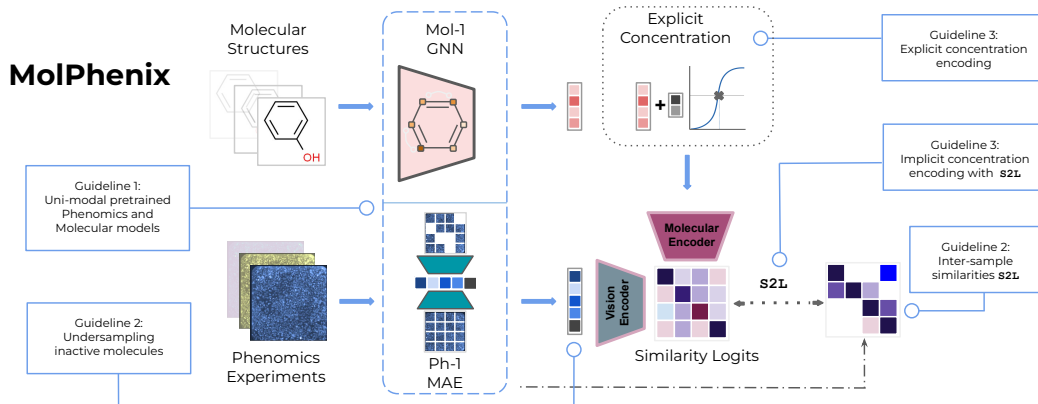


Figure 1: Illustration of proposed guidelines when incorporated in our *MolPhenix* contrastive phenomolecular retrieval framework. We address challenges by utilizing uni-modal pretrained MAE & MPNN models, inter-sample weighting with a dosage aware S2L loss, undersampling inactive molecules, and encoding molecular concentration.

Following these principles, we build *MolPhenix*, a multi-modal molecular phenomics model addressing contrastive phenomolecular retrieval (Figure 1). MolPhenix demonstrates large and consistent improvements in the presence of batch effects, generalizing across different concentrations, molecules, and activity thresholds. Additionally, MolPhenix outperforms baseline methods in zero-shot setting, achieving 77.33% top-1% retrieval accuracies on active molecules, which corresponds to a **8.1** $\times$  improvement over the previous state-of-the-art (SOTA) Sanchez-Fernandez et al. (2023).

## 2 RELATED WORK

**Uni-modality Pretraining:** Self-supervised methods have demonstrated success across a variety of domains such as computer vision, natural language processing and molecular representations [Balestriero et al. \(2023\)](#); [Radford et al. \(2019\)](#); [Zaidi et al. \(2022\)](#). In vision, contrastive methods have been used to minimize distance in the model’s latent space of two views of the same sample [Chen et al. \(2020\)](#); [Sohn \(2016\)](#); [Grill et al. \(2020\)](#); [He et al. \(2020\)](#). Reconstruction objectives have also permeated computer vision, such as masked autoencoders (MAE). MAEs typically utilize vision transformers to partition the image into learnable tokens and reconstruct masked patches [He et al. \(2022\)](#); [Feichtenhofer et al. \(2022\)](#); [Cao et al. \(2022\)](#); [Dosovitskiy et al. \(2020\)](#). These methods have been extended to microscopy experimental data designed to capture cell morphology [Xie et al. \(2023\)](#); [Kraus et al. \(2024\)](#). Ph-1 utilizes a masked autoencoder with a ViT-L/8+ architecture and a custom Fourier domain reconstruction loss, yielding informative representations of phenomic experiments [Kraus et al. \(2024\)](#); [Dehghani et al. \(2023\)](#). From a representational perspective, Graph Neural Networks (GNN) have been used to predict molecular properties by reasoning over graph structures. A combination of reconstruction and supervised objectives have led to models generalizing to a diverse range of prediction tasks [Méndez-Lucio et al. \(2022\)](#); [Zhou et al. \(2023\)](#); [Sypetkowski et al. \(2024\)](#); [Rong et al. \(2020\)](#). Our work leverages uni-modal foundation models, which are used to generate embeddings of phenomic images and molecular graphs.

**Multi-Modal Objectives:** Multi-modal models combine samples from two or more domains, to learn rich representations and demonstrate flexible ways to predict sample properties [Radford et al. \(2021\)](#); [Alayrac et al. \(2022\)](#); [Huang et al. \(2023\)](#). Contrastive methods minimize distances between paired samples, traditionally in text-image domains. However, training these models is computationally expensive, requiring large datasets. Multiple contributions have allowed for a reduction in compute and data budgets by an order of magnitude. In *LiT*, the authors demonstrate that utilizing uni-modal pretrained models for one or both modalities matches zero-shot performance with an order of magnitude fewer paired examples seen [Zhai et al. \(2022\)](#). [Zhai et al. \(2023\)](#) demonstrate that by replacing the softmax operation over cosine similarities with an element wise sigmoid loss, allows contrastive learners to improve performance under label noise regime [Zhai et al. \(2023\)](#). By using a uni-modal pre-trained modal to calculate similarities between samples from one of the modalities, [Srinivasa et al. \(2023\)](#) have demonstrated improved performance on zero-shot evaluation [Srinivasa et al. \(2023\)](#). In our work, we build along these directions in molecular phenomic multi-modal training.

**Molecular-Phenomic Contrastive Learning:** Prior works in contrastive phenomic retrieval have utilized the InfoNCE objective as a pre-training technique to construct uni-modal representations [Oord et al. \(2018\)](#). Recent methods have attempted to improve retrieval by using the InfoLOOB objective [Poole et al. \(2019\)](#). Specifically, CLOOME utilizes the InfoLOOB loss with hopfield networks for zero-shot retrieval on unseen data samples [Ramsauer et al. \(2020\)](#); [Sanchez-Fernandez et al. \(2023\)](#). Our work is parallel to the above directions, demonstrating a significant increase in molecular-phenomic retrieval by building on algorithmic improvements from the multi-modality literature.

## 3 METHODOLOGY

In this section, we explain key challenges facing phenomolecular retrieval and provide guidelines that are key methodological improvements behind the success of MolPhenix 1.

**Preliminaries:** Our setting studies the problem of learning multi-modal representations of molecules and phenomic experiments of treated cells [Sanchez-Fernandez et al. \(2023\)](#). The aim of this work is to learn a joint latent space which maps phenomic experiments of treated cells and the corresponding molecular perturbations into the same latent space. We consider a set of lab experiments  $\mathcal{E}$  defined as the tuple  $(\mathbf{X}, \mathbf{M}, \mathbf{C}, \Psi)$ . Each experiment  $\epsilon \in \mathcal{E}$  consists of data samples  $\mathbf{x}_i \in \mathbf{X}$  (such as images) and perturbations  $\mathbf{m}_i \in \mathbf{M}$  (such as molecules) which are obtained at a specific dosage concentration  $\mathbf{c}_i \in \mathbf{C}$ , while  $\psi \in \Psi$  denotes molecular activity threshold.

Figure 2 describes the problem of contrastive phenomolecular retrieval, where for a single image  $\mathbf{x}_i$ , the challenge consists of identifying the matching perturbation,  $\mathbf{m}_i$ , and concentration,  $\mathbf{c}_i$ , used to induce morphological effects. This can be accomplished in a zero-shot way by generating embeddings for  $(\mathbf{m}_1, \mathbf{c}_1), \dots, (\mathbf{m}_j, \mathbf{c}_j)$  and  $\mathbf{x}_i$  using functions  $f_{\theta_m}(\mathbf{m}, \mathbf{c})$ ,  $f_{\theta_x}(\mathbf{x})$  which map samples into  $\mathbb{R}^d$ . Then, by defining a similarity metric between generated embeddings  $\mathbf{z}_{x_i}$  and  $\mathbf{z}_{m_i}$ ,  $f_{sim}$ , we can rank  $(\mathbf{m}_1, \mathbf{c}_1) \dots (\mathbf{m}_j, \mathbf{c}_j)$  based on computed similarities. An effective solution to the contrastive phenomolecular retrieval problem would learn  $f_{\theta_m}(\mathbf{m}, \mathbf{c})$  and  $f_{\theta_x}(\mathbf{x})$  that results in consistently high retrieval rates of  $(\mathbf{m}_i, \mathbf{c}_i)$  used to perturb  $\mathbf{x}_i$ .

In practice, the image embeddings are generated using a phenomics microscopy foundation MAE model Kraus et al. (2024); He et al. (2022). We use phenomic embeddings to marginalize batch effects, infer inter-sample similarities, and undersample inactive molecules. Activity is determined using consistency of replicate measurements for a given perturbation. For each sample, a  $p$  value cutoff  $\psi \in \Psi$  is used to quantify molecular activity. Only molecules below the  $p$  value cutoff  $\psi$  are considered active.

Prior methods in multi-modal contrastive learning utilize the InfoNCE loss, and variants thereof Oord et al. (2018) to maximize the joint likelihood of  $\mathbf{x}_i$  and  $\mathbf{m}_i$ . Given a set of  $N \times N$  random samples  $(\mathbf{x}_1, \mathbf{m}_1, \mathbf{c}_1), \dots, (\mathbf{x}_N, \mathbf{m}_N, \mathbf{c}_N)$  containing  $N$  positive samples at  $k^{\text{th}}$  index and  $(N - 1) \times N$  negative samples, optimizing Equation 1 maximizes the likelihood of positive pairs while minimizing the likelihood of negative pairs:

$$\mathcal{L}_{\text{InfoNCE}} = -\frac{1}{N} \sum_{i=1}^N \left[ \log \frac{\exp(\langle \mathbf{z}_{x_i}, \mathbf{z}_{m_i} \rangle / \tau)}{\sum_{k=1}^N \exp(\langle \mathbf{z}_{x_i}, \mathbf{z}_{m_k} \rangle / \tau)} + \log \frac{\exp(\langle \mathbf{z}_{x_i}, \mathbf{z}_{m_i} \rangle / \tau)}{\sum_{k=1}^N \exp(\langle \mathbf{z}_{m_i}, \mathbf{z}_{x_k} \rangle / \tau)} \right]. \quad (1)$$

Where  $\mathbf{z}_x, \mathbf{z}_m$  correspond to phenomics and molecular embeddings respectively,  $\tau$  is softmax temperature, and  $\langle \cdot \rangle$  corresponds to cosine similarity.

#### CHALLENGE 1: PHENOMIC PRETRAINING AND GENERALIZATION

We find that using a phenomics foundation model to embed microscopy images allows for mitigation of batch effects, reduces the required number of paired data points, and improves generalization in the process. While CLIP, a hallmark model in the field of text-image multi-modality, was trained on 400 million curated paired data points, there is an order of magnitude fewer paired molecular-phenomic molecule samples Radford et al. (2021). Cost and systematic pre-processing of data make large scale data generation efforts challenging, and resulting data is affected by experimental batch effects. **Batch effects** induce noise in the latent space as a result of random perturbations in the experimental process, while biologically meaningful variation remains unchanged Parker & Leek (2012); Sonesson et al. (2014). Limited dataset sizes and batch effects make it challenging for contrastive learners to capture molecular features affecting cell morphology, yielding low retrieval rates Sanchez-Fernandez et al. (2023).

We address data availability and generalization challenges by utilizing representations from a large **uni-modal pre-trained phenomic model**,  $\theta_{\text{ph}}$ , trained to capture representations of cellular morphology.  $\theta_{\text{ph}}$  is pretrained on microscopy images using a Fourier modified MAE objective, utilizing the ViT-L/8 architecture with methodology similar to Kraus et al. (2024) He et al. (2022); Dosovitskiy et al. (2020); Kraus et al. (2024). For simplicity in future

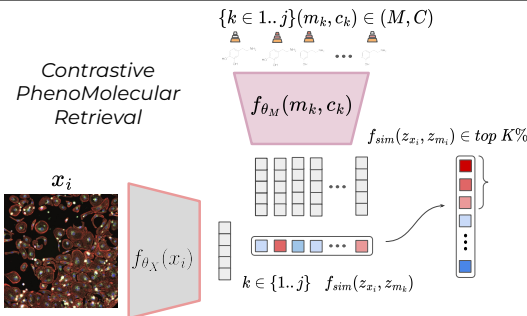


Figure 2: Illustration of the contrastive phenomolecular retrieval challenge. Image  $\mathbf{x}_i$  and a set of molecules and corresponding concentrations  $(\mathbf{m}_k, \mathbf{c}_k)$  get mapped into a  $\mathbb{R}^d$  latent space. Their similarities get computed with  $f_{sim}$  and ranked to evaluate whether the paired perturbation appears in the top  $K\%$ .


sections, we refer to this model as *Ph-1*. This pretrained model allows a drastic reduction in the required number of paired multi-modal samples [Zhai et al. \(2022\)](#). In addition, using phenomic representations alleviates the challenge of batch effects by averaging samples,  $\mathbf{z}_{x_i}$ , generated with the same perturbation  $\mathbf{m}_i$  over multiple lab experiments  $\epsilon_i$ . Averaging model representations  $\frac{1}{N} \sum_{i \in N} \mathbf{z}_{x_i}$  allows marginalizing batch effect induced by individual experiments.

**Guideline 1** Utilizing pre-trained uni-modal encoder,  $\theta_{Ph}$ , can be used to reduce the number of paired data-points compared to training  $\theta$  without prior optimization. In addition, averaging phenomic embeddings  $\mathbf{z}_x$  from matched perturbations can alleviate batch effects.

To reason over molecular structures, we make use of features learned from GNNs trained on molecular property prediction [Masters et al. \(2023\)](#). We utilize a pretrained MPNN foundational model up to the order of 1B parameters for extracting molecular representations following a similar procedure to [Sypetkowski et al. \(2024\)](#) [Sypetkowski et al. \(2024\)](#). We refer to this model as *Mol-1*.

## CHALLENGE 2: INACTIVE MOLECULAR PERTURBATIONS

The phenomics-molecular data collection process can result in pairing of molecular structures with unperturbed cells in cases where the molecule has no effect on cell morphology (Figure 3)



Since the morphological effects observed in cell  $x_i$  is conditioned on the perturbation, in the absence of a molecular effect  $P(x_i | x_i^0, c_i, m_i) \sim P(x_i | x_i^0)$ . In these samples, phenomic data will be independent, from paired molecular data, which results in misannotation under the assumption of data-pairs having an underlying biological relationship. We demonstrate how utilizing *Ph-1* to undersample inactive molecules and learn continuous similarities between samples can alleviate this challenge.

Figure 3: Data generation process of a phenomic experiment on cells  $x_i$  with molecular perturbations  $m_i$  and concentrations  $c_i$ .

To **undersample inactive molecules**, we extract the embeddings from *Ph-1* and calculate the relative activity of each perturbation  $(m_i, c_i) \in (M, C)$ . This is done using the rank of cosine similarities between technical replicates produced for a molecular perturbation against a null distribution. The null distribution is established by calculating cosine similarities from random pairs of *Ph-1* embeddings generated with perturbation  $(m_j, c_j)$ ,  $(m_k, c_k)$ . Hence, we can compute a p-value and filter out samples likely to belong to the null distribution with an arbitrary threshold  $\psi$ .

In addition, by utilizing an inter-sample aware **S2L training objective**, the model can learn similarities between inactive molecules. S2L is grounded in previous work which demonstrates improved robustness to label noise (SigLip) and learnable inter-sample associations (CWCL) [Zhai et al. \(2023\)](#); [Srinivasa et al. \(2023\)](#). Continuous Weighted Contrastive Loss (CWCL) provides better multi-modal alignment using a uni-modal pretrained model to suggest sample distances, relaxing the negative equidistant assumption present in InfoNCE [Srinivasa et al. \(2023\)](#):

$$\mathcal{L}_{CWCL, \mathcal{M} \rightarrow \mathcal{X}} = -\frac{1}{N} \sum_{i=1}^N \left[ \frac{1}{\sum_{j=1}^N \mathbf{w}_{i,j}^{\mathcal{X}}} \sum_{i=1}^N \mathbf{w}_{i,j}^{\mathcal{X}} \log \frac{\exp(\langle \mathbf{z}_{x_i}, \mathbf{z}_{m_j} \rangle / \tau)}{\sum_{k=1}^N \exp(\langle \mathbf{z}_{x_j}, \mathbf{z}_{m_k} \rangle / \tau)} \right]. \quad (2)$$

CWCL weights logits with a continuous measure of similarity  $\mathbf{w}^{\mathcal{X}}$ , resulting in better alignment of embeddings  $\mathbf{z}_{x_i}$  and  $\mathbf{z}_{m_j}$  across modalities. In equation 2,  $\mathbf{w}^{\mathcal{X}}$  is computed using a within modality similarity function such as  $\mathbf{w}_{i,j}^{\mathcal{X}} = \langle \mathbf{z}_{x_i}, \mathbf{z}_{x_j} \rangle / 2 + 0.5$ . Note, the above formula



is used only for mapping samples from modality  $\mathcal{M}$  to  $\mathcal{X}$  for which a pre-trained model  $\theta_{\text{ph}}$  is available.

Another work, SigLIP, demonstrates robustness to label noise and reduces computational requirements during contrastive training Zhai et al. (2023). It does so by avoiding computation of a softmax over the entire set of in-batch samples, instead relying on element-wise sigmoid operation:

$$\mathcal{L}_{\text{SigLIP}} = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^N \left[ \log \frac{1}{1 + \exp(\mathbf{y}_{i,j}(-\alpha \langle \mathbf{z}_{\mathbf{x}_i}, \mathbf{z}_{\mathbf{m}_j} \rangle + b))} \right]. \quad (3)$$

In equation 3,  $\alpha$  and  $b$  are learned, calibrating the model confidence conditioned on the ratio of positive to negative pairs.  $\mathbf{y}_{i,j}$  is set to 1 if  $i = j$  and -1 otherwise.

Inspired by prior works, we introduce S2L for molecular representation learning, which leverages inter-sample similarities and robustness to label noise to mitigate weak or inactive perturbations.

$$\mathcal{L}_{\text{S2L}} = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^N \log \left[ \frac{\mathbf{w}_{i,j}^{\mathcal{X}}}{1 + \exp(-\alpha \mathbf{z}_{\mathbf{x}_i} \cdot \mathbf{z}_{\mathbf{m}_j} + b)} + \frac{(1 - \mathbf{w}_{i,j}^{\mathcal{X}})}{1 + \exp(\alpha \mathbf{z}_{\mathbf{x}_i} \cdot \mathbf{z}_{\mathbf{m}_j} + b)} \right]. \quad (4)$$

In the equation above,  $\mathbf{z}_{\mathbf{x}_i}$  and  $\mathbf{z}_{\mathbf{m}_j}$  correspond to latent representations of images and molecules, respectively.  $\alpha$  and  $b$  correspond to learnable temperature and bias parameters for the calibrated sigmoid function.  $\mathbf{w}_{i,j}^{\mathcal{X}}$  is an inter-sample similarity function computed from images using the pretrained model  $\theta_{\text{ph}}$ . To compute  $\mathbf{w}_{i,j}^{\mathcal{X}}$ , we use the arctangent of L2 distance instead of a modified cosine similarity, as was the case for Equation 2 (more details in Appendix C.3). Intuitively, S2L can be thought of as shifting from a multi-class classification to a soft multi-label problem. In our problem setting, the labels are continuous and determined by sample similarity in the phenomics space.

**Guideline 2** When training a molecular-phenomic model, mitigating the effect of inactive molecules in training data distribution can be carried out by undersampling inactive molecules and using an inter-sample similarity aware, S2L loss (equation 4).

### CHALLENGE 3: VARIABLE CONCENTRATIONS

Perturbation effect on a cell is determined by both molecular structure and corresponding concentration Walmsley & Billinton (2011). A model capturing molecular impact on cell morphology must be able to generalize across different doses, since variable concentrations can correspond to different data distributions.

We note that providing concentrations  $c_i$  as input to the model would benefit performance, as this would indicate the magnitude of molecular impact. However, we find that simply concatenating concentrations does not result in effective training due to its compressed dynamic range. To that end, we add concentration information in two separate ways: *implicit* and *explicit* formulations.

We add **implicit concentration** as molecular perturbation classes by using the S2L loss (equation 4) to treat perturbation  $\mathbf{m}_i$  with concentrations  $c_i$  and  $c_j$  as distinct classes. This pushes samples apart in the latent space proportionally to similarities between phenomic experiments.

We add **explicit concentration**  $c_i$  by passing it to the molecular encoder. We explore different formulation for dosage concentrations,  $\mathbf{f}'(c_i)$ , where  $\mathbf{f}'$  maps  $c_i \rightarrow \mathbb{R}$ . Encoded representations  $\mathbf{f}'(c_i)$  are concatenated at the initial layer of the model. We find simple functional encodings  $\mathbf{f}'$  (such as one-hot and logarithm) to work well in practice.

**Guideline 3** When training a molecular-phenomic model, conditioning on an (implicit and explicit) representation of concentration  $\mathbf{f}'(\mathbf{c}_i)$  aids in capturing molecular impacts on cell morphology and improves generalization to previously unseen molecules and concentrations.

## 4 EXPERIMENTAL SETUP

In this section, we describe evaluation datasets used, and descriptions of the underlying data modalities. To assess phenomolecular retrieval, we use 1% recall metric unless stated otherwise, as it allows direct comparison between datasets with different number of samples. Additional implementation and evaluation details can be found in Appendix C.

**Datasets:** Our training dataset consists of fluorescent microscopy images paired with molecular structures and concentrations, which are used as perturbants. We assess models’ phenomolecular retrieval capabilities on three datasets of escalating generalization complexity. First dataset, consisting of unseen microscopy images and molecules present in the training dataset. Second, a dataset consisting of previously unseen phenomics experiments and molecules split by the corresponding molecular scaffold. Finally, we evaluate on an open source dataset with a different data generating distribution [Fay et al. \(2023\)](#). In the case of the latter two datasets, the model is required to perform zero-shot classification, as it has no access to those molecules in the training data. This requires the model to reason over molecular graphs to identify structures inducing corresponding cellular morphology changes. Using methodology described in guideline 2 we report retrieval results for all molecules as well as on an active subset. Finally, all datasets are comprised of molecular structures at multiple concentrations (.01, .1, 1.0, 10, etc.) Additional details regarding the datasets can be found in Appendix B.

**Modality Representations:** In our evaluations, we consider different representations for molecular perturbations and phenomic experiments and quantitatively evaluate their impact.

- **Images:** Image encoders utilize 6-channel fluorescent microscopy images of cells representing phenomic experiments. Images are  $2048 \times 2048$  pixels, capturing cellular morphology changes post molecular perturbation. We downscale each image to  $256 \times 256$  using block mean downsampling.
- **ph-1:** We characterize phenomic experiments by embedding high resolution microscopy images in the latent space of a phenomics model  $\theta_{Ph}$  as described in guideline 1.
- **Fingerprints:** Molecular fingerprints utilize RDKit [Landrum et al. \(2013\)](#), MACCS [Kuwahara & Gao \(2021\)](#) and MORGAN3 [Rogers & Hahn \(2010\)](#) bit coding, which represent binary presence of molecular substructures. Additional information such as atomic identity, atomic radius and torsional angles are included in the fingerprint representations.
- **Mol-1:** We generate molecular representations from a large pretrained GNN. Specifically, we obtain molecular embeddings from a 1B parameter MPNN [Masters et al. \(2023\)](#).

## 5 RESULTS AND DISCUSSION

To evaluate the effectiveness of guidelines 1, 2, and 3 we carry out evaluation in two different settings: (1) cumulative concentrations, and (2) held-out concentrations, testing the models’ ability to generalize to new molecular doses. Finally, we perform comprehensive ablations testing model performance with varying data, model, and optimization parameters. The comprehensive set of results can be found in Tables 8, 10, 9 and 11.

### 5.1 EVALUATION ON CUMULATIVE CONCENTRATIONS



We demonstrate improvements in phenomolecular recall due to usage of a phenomics pre-trained foundation model, identify that MolPhenix set of design choices results in higher final performance, and more data efficient learning. Figure 4 demonstrates recall accuracy on all molecules and an active subset for CLOOME and MolPhenix models, as a function of training samples seen.

We observe a large performance gap between models trained on ph-1 embeddings as opposed to images, emphasizing the utility of using a pre-trained encoder for microscopy images (Table 1). We note that provision of ph-1 (CLOOME-Ph-1 Vs CLOOME-Images) significantly improves both active and all molecule retrieval by  $5.69\times$  and,  $4.75\times$  respectively (Table 1). Furthermore, we identify that while all molecules retrieval stagnates throughout training, the performance on an active subset keeps improving, underscoring the importance of identification of the active subset. Finally, we compare CLOOME and MolPhenix trained using ph-1 embeddings and find there is a consistent retrieval performance gap, throughout training, with a  $1.26\times$  final improvement (Figure 4, Table 1). Compared to CLOOME [Sanchez-Fernandez et al. \(2023\)](#) trained directly on images, MolPhenix achieves an average improvement of  $8.78\times$  on active molecules on the unseen dataset. These results verify the effectiveness of guideline 1 in accelerating training, and the importance of guidelines 2 and 3 in recall improvements over CLOOME.




Figure 4: Comparison of training phenomic encoder from scratch and utilizing pre-trained Ph-1 unseen dataset. X-axis plotted on logarithmic scale.

We evaluate the impact of different loss objectives on the proposed MolPhenix training framework. Table 2 presents top-1% retrieval accuracy across different contrastive losses utilized to train molecular-phenomics encoders on cumulative concentrations. Compared to prior methods, the proposed S2L loss demonstrates improved retrieval rates in cumulative concentration setting. Label noise and inter-sample similarity aware losses such as CWCL and SigLip also demonstrate improved performance. Effectiveness of S2L can be attributed to smoothed inter-sample similarities and implicit concentration information.

Finally, in Table 3, we observe recall improvements when considering both molecular structures and concentration. We note the importance of addition of implicit concentration, further confirming the importance of considering molecular effects at different concentrations as different classes. Explicitly encoding molecular concentration with one-hot, logarithm and sigmoid yields improved recall performance, with one-hot perform the best in a cumulative concentration setting, These findings verify the efficacy of implicit and explicit concentration encoding outlined in guideline 3.

Table 1: Impact of pre-trained ph-1 and Mol-1 on CLOOME and MolPhenix for a matched number of seen samples (Top), where we observe an  $8.1\times$  improvement of MolPhenix over the CLOOME baseline for active unseen molecules. SOTA results trained with a higher number of steps by utilizing the best hyperparameters (Bottom \*). We note that MolPhenix’s main components such as S2L and embeddings averaging relies on having a pre-trained uni-modal phenomics model and so cannot be fully adopted to being trained on images.

Method	Modality	Active Molecules			All Molecules		
		Unseen Im.	Unseen Im. + Mol.	Unseen Dataset	Unseen Im.	Unseen Im. + Mol.	Unseen Dataset
CLOOME	Images & Multi-FPS	.0756 $\pm$ .0042	.0787 $\pm$ .0065	.0528 $\pm$ .0057	.0547 $\pm$ .0028	.0661 $\pm$ .0020	.0223 $\pm$ .0014
CLOOME	ph-1 & Multi-FPS	.4659 $\pm$ .0042	.5057 $\pm$ .0014	.2065 $\pm$ .0146	.3009 $\pm$ .0053	.2474 $\pm$ .0013	.1337 $\pm$ .0045
MolPhenix	ph-1 & Multi-FPS	<b>.7807 <math>\pm</math> .0025</b>	.6365 $\pm$ .0014	.3545 $\pm$ .0097	<b>.5253 <math>\pm</math> .0029</b>	<b>.3655 <math>\pm</math> .0017</b>	.2163 $\pm$ .0021
MolPhenix	ph-1 & Mol-1	.7646 $\pm$ .0014	<b>.6387 <math>\pm</math> .0056</b>	<b>.4160 <math>\pm</math> .0016</b>	.5012 $\pm$ .0002	.3511 $\pm$ .0004	<b>.2508 <math>\pm</math> .0026</b>
MolPhenix*	ph-1 & Mol-1	<b>.9689 <math>\pm</math> .0017</b>	<b>.7733 <math>\pm</math> .0036</b>	<b>.5860 <math>\pm</math> .0082</b>	<b>.5583 <math>\pm</math> .0007</b>	<b>.3824 <math>\pm</math> .0016</b>	<b>.2809 <math>\pm</math> .0060</b>

Table 2: Top-1% recall accuracy with use of the proposed MolPhenix guidelines, such as ph-1 and embedding averaging. We omit explicit concentration from this experiment.

Loss	Active Molecules			All Molecules		
	Unseen Images	Unseen Im. + Mol.	Unseen Dataset	Unseen Images	Unseen Im. + Mol.	Unseen Dataset
CLIP	.3373 $\pm$ .0043	.4228 $\pm$ .0008	.1514 $\pm$ .0038	.1761 $\pm$ .0043	.1867 $\pm$ .0022	.0734 $\pm$ .0022
Hopfield-CLIP	.2578 $\pm$ .0042	.3559 $\pm$ .0042	.1256 $\pm$ .0092	.1531 $\pm$ .0046	.1709 $\pm$ .0029	.0673 $\pm$ .0020
InfoLOOB	.3351 $\pm$ .0011	.4206 $\pm$ .0031	.1563 $\pm$ .0028	.1746 $\pm$ .0003	.1860 $\pm$ .0029	.0745 $\pm$ .0019
CLOOME	.3572 $\pm$ .0026	.4348 $\pm$ .0039	.1658 $\pm$ .0063	.1968 $\pm$ .0029	.2005 $\pm$ .0026	.0911 $\pm$ .0022
DCL	.6363 $\pm$ .0025	.6177 $\pm$ .0047	.3184 $\pm$ .0087	.3277 $\pm$ .0047	.2562 $\pm$ .0008	.1364 $\pm$ .0067
CWCL	.7091 $\pm$ .0045	.6529 $\pm$ .0020	.3556 $\pm$ .0094	.3635 $\pm$ .0064	.2696 $\pm$ .0019	.1526 $\pm$ .0058
SigLip	.7763 $\pm$ .0045	.6401 $\pm$ .0065	.3396 $\pm$ .0042	.3729 $\pm$ .0039	.2544 $\pm$ .0014	.1470 $\pm$ .0038
S2L (ours)	<b>.9097 <math>\pm</math> .0020</b>	<b>.6759 <math>\pm</math> .0012</b>	<b>.4181 <math>\pm</math> .0012</b>	<b>.4688 <math>\pm</math> .0009</b>	<b>.2852 <math>\pm</math> .0001</b>	<b>.1838 <math>\pm</math> .0007</b>

Table 3: Top-1% recall accuracy across different concentration encoding choices with use of the proposed MolPhenix guidelines, such as ph-1 and embedding averaging.

Implicit Concentration	Explicit Concentration	Active Molecules			All Molecules		
		Unseen Im.	Unseen Im. + Mol.	Unseen Dataset	Unseen Im.	Unseen Im. + Mol.	Unseen Dataset
$\times$	$\times$	.7350 $\pm$ .0071	.6509 $\pm$ .0104	.3333 $\pm$ .0004	.3610 $\pm$ .0025	.2668 $\pm$ .0034	.1532 $\pm$ .0007
$\checkmark$	$\times$	.9097 $\pm$ .0020	.6759 $\pm$ .0012	.4181 $\pm$ .0012	.4688 $\pm$ .0009	.2852 $\pm$ .0001	.1838 $\pm$ .0007
$\checkmark$	sigmoid	.9423 $\pm$ .0011	.7155 $\pm$ .0016	.4573 $\pm$ .0022	.5071 $\pm$ .0024	.3441 $\pm$ .0026	.2144 $\pm$ .0026
$\checkmark$	logarithm	.9426 $\pm$ .0066	.7451 $\pm$ .0050	.4727 $\pm$ .0056	.5183 $\pm$ .0027	.3700 $\pm$ .0036	.2275 $\pm$ .0032
$\checkmark$	one-hot	<b>.9430 <math>\pm</math> .0029</b>	<b>.7490 <math>\pm</math> .0052</b>	<b>.4850 <math>\pm</math> .0020</b>	<b>.5433 <math>\pm</math> .0030</b>	<b>.3819 <math>\pm</math> .0032</b>	<b>.2384 <math>\pm</math> .0049</b>

Table 4: Top-1% recall accuracy of different loss objectives while using the proposed MolPhenix guidelines, such as ph-1 and embedding averaging.

Loss	Unseen Im.	Unseen Im. + Mol.	Unseen Dataset
CLIP	.2109	.2425	.1519
Hopfield-CLIP	.1581	.2034	.1198
InfoLOOB	.2122	.2496	.1501
CLOOME	.2164	.2461	.1479
DCL	.4717	.4027	.2841
CWCL	.5731	.4403	.3232
SigLip	.5718	.4217	.3021
S2L (ours)	<b>.8334</b>	<b>.4615</b>	<b>.3792</b>

Results are averaged across experiments for each dropped concentration, and across three seeds. Recall is reported for active molecules, while the results for all molecules can be found in Table 11.

## 5.2 EVALUATION ON HELD-OUT CONCENTRATIONS

Next, we evaluate recall on held-out concentrations to obtain a measure of generalization performance. This evaluation allows us to capture the utility of our models for prediction of unseen concentrations, hence resembling *in-silico* testing. We omit concentrations from the training set and evaluate recall at the excluded data, where we observe a drop in retrieval performance for unseen concentrations. Similar to cumulative concentration results, we find that using S2L improves recall over other losses and outperforms CLOOME by up to 126% (Table 4). While one-hot encoding exhibits significant improvements in cumulative concentrations, its expressivity on unseen concentrations is limited (Table 5) and sigmoid encoding provides a sufficient representation of concentration.

## 5.3 ABLATION STUDIES

We assess the importance of our design decisions by conducting an ablation study over our proposed guidelines. Figure 5 presents the variation of top-1% recall accuracy across

key components such as cutoff  $p$  value, fingerprint type, and embedding averaging. We observe that employing a lower cutoff  $p$  value yields improved generalization for unseen dataset, while employing a higher cutoff appears to be optimal for unseen images + unseen molecules. For molecular structure representations, we find that using embeddings from the large pretrained MPNN graph based model (e.g., Mol-1) surpasses traditional fingerprints. Finally, utilization of embedding averaging demonstrates improved recall. More ablations over model size, projector dimension, and batch size can be found in appendix D.4.




Figure 5: Ablations of top-1 % recall accuracy with **(bottom-left)** cutoff  $p$  value, **(bottom-center)** fingerprint type, and **(bottom-right)** embedding averaging.

## 6 CONCLUSION

In this work, we investigate the problem of *contrastive phenomolecular retrieval* by constructing a joint multi-modal embedding of phenomic experiments and molecular structures. We identify a set of challenges afflicting molecular-phenomic training and proposed a set of guidelines for improving retrieval and generalization. Empirically, we observed that contrastive learners demonstrate higher retrieval rates when using representations from a high-capacity uni-modal pretrained model. Use of inter-sample similarities with a label noise resistant loss such as S2L allows us to tackle the challenge of inactive molecules. Finally, adding implicit and explicit concentrations allows models to generalize to previously unseen concentrations. MolPhenix demonstrates an  $8.1\times$  improvement in zero shot retrieval of active molecules over the previous state-of-the-art, reaching 77.33% in top-1% accuracy. In addition, we conduct a preliminary investigation on MolPhenix’s ability to uncover biologically meaningful properties D.2D.1. We expect a wide range of applications for MolPhenix, particularly in drug discovery. While there’s a remote chance of misuse for developing chemical weapons, such harm is unlikely, with our primary focus remaining on healthcare improvement.

**Limitations and Future Work:** While our study covers challenges in phenomolecular recall, we leave three research directions for future work. (1) Future investigations could consider studying additional modalities such as text, genetic perturbations and chemical multi-compound interventions. (2) While we propose and evaluate our guidelines on previously conducted phenomic experiments, we note that a rigorous evaluation would evaluate model predictions in a wet-lab setting. (3) In addition, our work makes the assumption that the initial unperturbed cell state  $x_i^0$  can be marginalized by utilizing a single cell line with an unperturbed genetic background. Future works can relax this assumption, aiming to capture innate intercellular variation.

## 7 ACKNOWLEDGEMENTS

We thank the broader Valence Labs team and Recursion Pharmaceuticals for support on the project. We thank Berton Earnshaw, Jason Hartford and Emmanuel Noutahi for providing valuable feedback.

## REFERENCES

- Jean-Baptiste Alayrac, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, Arthur Mensch, Katie Millican, Malcolm Reynolds, Roman Ring, Eliza Rutherford, Serkan Cabi, Tengda Han, Zhitao Gong, Sina Samangooei, Marianne Monteiro, Jacob Menick, Sebastian Borgeaud, Andrew Brock, Aida Nematzadeh, Sahand Sharifzadeh, Mikolaj Binkowski, Ricardo Barreira, Oriol Vinyals, Andrew Zisserman, and Karen Simonyan. Flamingo: a visual language model for few-shot learning, 2022.
- Saleh Albelwi. Survey on self-supervised learning: auxiliary pretext tasks and contrastive learning methods in imaging. *Entropy*, 24(4):551, 2022.
- Randall Balestriero, Mark Ibrahim, Vlad Sobal, Ari Morcos, Shashank Shekhar, Tom Goldstein, Florian Bordes, Adrien Bardes, Gregoire Mialon, Yuandong Tian, Avi Schwarzschild, Andrew Gordon Wilson, Jonas Geiping, Quentin Garrido, Pierre Fernandez, Amir Bar, Hamed Pirsiavash, Yann LeCun, and Micah Goldblum. A cookbook of self-supervised learning, 2023.
- Christoph Bock, Paul Datlinger, Florence Chardon, Matthew A. Coelho, Matthew B. Dong, Keith A. Lawson, Tian Lu, Laetitia Maroc, Thomas M. Norman, Bicna Song, Geoff Stanley, Sidi Chen, Mathew Garnett, Wei Li, Jason Moffat, Lei S. Qi, Rebecca S. Shapiro, Jay Shendure, Jonathan S. Weissman, and Xiaowei Zhuang. High-content crispr screening. *Nature Reviews Methods Primers*, 2(1), February 2022. ISSN 2662-8449. doi: 10.1038/s43586-021-00093-4. URL <http://dx.doi.org/10.1038/s43586-021-00093-4>.
- Regine S. Bohacek, Colin McMartin, and Wayne C. Guida. The art and practice of structure-based drug design: A molecular modeling perspective. *Medicinal Research Reviews*, 16(1):3-50, January 1996. ISSN 1098-1128. doi: 10.1002/(sici)1098-1128(199601)16:1<3::aid-med1>3.0.co;2-6. URL [http://dx.doi.org/10.1002/\(SICI\)1098-1128\(199601\)16:1<3::AID-MED1>3.0.CO;2-6](http://dx.doi.org/10.1002/(SICI)1098-1128(199601)16:1<3::AID-MED1>3.0.CO;2-6).
- Michael Boutros, Florian Heigwer, and Christina Laufer. Microscopy-based high-content screening. *Cell*, 163(6):1314-1325, 2015.
- Mark-Anthony Bray, Shantanu Singh, Han Han, Chadwick T Davis, Blake Borgeson, Cathy Hartland, Maria Kost-Alimova, Sigrun M Gustafsdottir, Christopher C Gibson, and Anne E Carpenter. Cell painting, a high-content image-based assay for morphological profiling using multiplexed fluorescent dyes. *Nature protocols*, 11(9):1757-1774, 2016.
- Shuhao Cao, Peng Xu, and David A Clifton. How to understand masked autoencoders. *arXiv preprint arXiv:2202.03670*, 2022.
- Srinivas Niranj Chandrasekaran, Jeanelle Ackerman, Eric Alix, D. Michael Ando, John Arevalo, Melissa Bennion, Nicolas Boisseau, Adriana Borowa, Justin D. Boyd, Laurent Brino, Patrick J. Byrne, Hugo Ceulemans, Carolyn Ch'ng, Beth A. Cimini, Djork-Arne Clevert, Nicole Deflaux, John G Doench, Thierry Dorval, Regis Doyonnas, Vincenzo Dragone, Ola Engkvist, Patrick W. Faloon, Briana Fritchman, Florian Fuchs, Sakshi Garg, Tamara J. Gilbert, David Glazer, David Gnuttt, Amy Goodale, Jeremy Grignard, Judith Guenther, Yu Han, Zahra Hanifelhoul, Santosh Hariharan, Desiree Hernandez, Shane R Horman, Gisela Hormel, Michael Huntley, Ilknur Icke, Makiyo Iida, Christina B. Jacob, Steffen Jaensch, Jawahar Khetan, Maria Kost-Alimova, Tomasz Krawiec, Daniel Kuhn, Charles-Hugues Lardeau, Amanda Lembke, Francis Lin, Kevin D. Little, Kenneth R. Lofstrom, Sofia Lotfi, David J. Logan, Yi Luo, Franck Madoux, Paula A. Marin Zapata, Brittany A. Marion, Glynn Martin, Nicola Jane McCarthy, Lewis Mervin, Lisa Miller, Haseeb Mohamed, Tiziana Monteverde, Elizabeth Mouchet, Barbara Nicke, Arnaud Ogier, Anne-Laure Ong, Marc Osterland, Magdalena Otrocka, Pieter J. Peeters, James Pilling, Stefan Prechtel, Chen Qian, Krzysztof Rataj, David E Root, Sylvie K. Sakata, Simon Scrace, Hajime Shimizu, David Simon, Peter Sommer, Craig Spruiell, Iffat Sumia, Susanne E Swalley, Hiroki Terauchi, Amandine Thibaudeau, Amy Unruh, Jelle Van de Waeter, Michiel Van Dyck, Carlo van Staden, Michał Warchoń, Erin Weisbart, Amélie Weiss, Nicolas Wiest-Daessle, Guy Williams, Shan Yu, Bolek Zapiec, Marek Żyła, Shantanu Singh, and

- Anne E. Carpenter. Jump cell painting dataset: morphological impact of 136,000 chemical and genetic perturbations. *bioRxiv*, 2023. doi: 10.1101/2023.03.23.534023. URL <https://www.biorxiv.org/content/early/2023/03/24/2023.03.23.534023>.
- Ting Chen, Simon Kornblith, Kevin Swersky, Mohammad Norouzi, and Geoffrey E Hinton. Big self-supervised models are strong semi-supervised learners. *Advances in neural information processing systems*, 33:22243–22255, 2020.
- Mostafa Dehghani, Josip Djolonga, Basil Mustafa, Piotr Padlewski, Jonathan Heek, Justin Gilmer, Andreas Peter Steiner, Mathilde Caron, Robert Geirhos, Ibrahim Alabdulmohsin, et al. Scaling vision transformers to 22 billion parameters. In *International Conference on Machine Learning*, pp. 7480–7512. PMLR, 2023.
- Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, et al. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020.
- Marta M Fay, Oren Kraus, Mason Victors, Lakshmanan Arumugam, Kamal Vuggumudi, John Urbanik, Kyle Hansen, Safiye Celik, Nico Cernek, Ganesh Jagannathan, et al. Rxrx3: Phenomics map of biology. *Biorxiv*, pp. 2023–02, 2023.
- Christoph Feichtenhofer, Yanghao Li, Kaiming He, et al. Masked autoencoders as spatiotemporal learners. *Advances in neural information processing systems*, 35:35946–35958, 2022.
- Andreas Fürst, Elisabeth Rumetshofer, Johannes Lehner, Viet T Tran, Fei Tang, Hubert Ramsauer, David Kreil, Michael Kopp, Günter Klambauer, Angela Bitto, et al. Cloob: Modern hopfield networks with infoloob outperform clip. *Advances in neural information processing systems*, 35:20450–20468, 2022.
- Jean-Bastien Grill, Florian Strub, Florent Altché, Corentin Tallec, Pierre H. Richemond, Elena Buchatskaya, Carl Doersch, Bernardo Avila Pires, Zhaohan Daniel Guo, Mohammad Gheshlaghi Azar, Bilal Piot, Koray Kavukcuoglu, Rémi Munos, and Michal Valko. Bootstrap your own latent: A new approach to self-supervised learning, 2020.
- Kaiming He, Haoqi Fan, Yuxin Wu, Saining Xie, and Ross Girshick. Momentum contrast for unsupervised visual representation learning, 2020.
- Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoencoders are scalable vision learners. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 16000–16009, 2022.
- Markus Hofmarcher, Elisabeth Rumetshofer, Djork-Arne Clevert, Sepp Hochreiter, and Gunter Klambauer. Accurate prediction of biological assays with high-throughput microscopy images and convolutional networks. *Journal of chemical information and modeling*, 59(3):1163–1171, 2019.
- Shaohan Huang, Li Dong, Wenhui Wang, Yaru Hao, Saksham Singhal, Shuming Ma, Tengchao Lv, Lei Cui, Owais Khan Mohammed, Barun Patra, Qiang Liu, Kriti Aggarwal, Zewen Chi, Johan Björck, Vishrav Chaudhary, Subhojit Som, Xia Song, and Furu Wei. Language is not all you need: Aligning perception with language models, 2023.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020.
- Craig Knox, Mike Wilson, Christen M Klinger, Mark Franklin, Eponine Oler, Alex Wilson, Allison Pon, Jordan Cox, Na Eun (Lucy) Chin, Seth A Strawbridge, Marysol Garcia-Patino, Ray Kruger, Aadhavya Sivakumaran, Selena Sanford, Rahil Doshi, Nitya Khetarpal, Omolola Fatokun, Daphnee Doucet, Ashley Zubkowski, Dorsa Yahya Rayat, Hayley Jackson, Karxena Harford, Afia Anjum, Mahi Zakir, Fei Wang, Siyang Tian, Brian Lee, Jaanus Liigand, Harrison Peters, Ruo Qi (Rachel) Wang, Tue Nguyen, Denise So,

- Matthew Sharp, Rodolfo da Silva, Cyrella Gabriel, Joshua Scantlebury, Marissa Jasin-ski, David Ackerman, Timothy Jewison, Tanvir Sajed, Vasuk Gautam, and David S Wishart. Drugbank 6.0: the drugbank knowledgebase for 2024. *Nucleic Acids Research*, 52(D1):D1265–D1275, November 2023. ISSN 1362-4962. doi: 10.1093/nar/gkad976. URL <http://dx.doi.org/10.1093/nar/gkad976>.
- Oren Kraus, Kian Kenyon-Dean, Saber Saberian, Maryam Fallah, Peter McLean, Jess Leung, Vasudev Sharma, Ayla Khan, Jia Balakrishnan, Safiye Celik, Dominique Beaini, Maciej Sypetkowski, Chi Vicky Cheng, Kristen Morse, Maureen Makes, Ben Mabey, and Berton Earnshaw. Masked autoencoders for microscopy are scalable learners of cellular biology, 2024.
- Oren Z Kraus, Ben T Grys, Jimmy Ba, Yolanda Chong, Brendan J Frey, Charles Boone, and Brenda J Andrews. Automated analysis of high-content microscopy data with deep learning. *Molecular systems biology*, 13(4):924, 2017.
- Hiroyuki Kuwahara and Xin Gao. Analysis of the effects of related fingerprints on molecular similarity using an eigenvalue entropy approach. *Journal of Cheminformatics*, 13:1–12, 2021.
- Greg Landrum et al. Rdkit: A software suite for cheminformatics, computational chemistry, and predictive modeling. *Greg Landrum*, 8(31.10):5281, 2013.
- Francois Lanusse, Liam Parker, Siavash Golkar, Miles Cranmer, Alberto Bietti, Michael Eickenberg, Geraud Krawezik, Michael McCabe, Ruben Ohana, Mariel Pettee, Bruno Regalado-Saint Blancard, Tiberiu Tesileanu, Kyunghyun Cho, and Shirley Ho. Astroclip: Cross-modal pre-training for astronomical foundation models, 2023.
- Jeffrey T Leek, Robert B Scharpf, Héctor Corrada Bravo, David Simcha, Benjamin Langmead, W Evan Johnson, Donald Geman, Keith Baggerly, and Rafael A Irizarry. Tackling the widespread and critical impact of batch effects in high-throughput data. *Nature Reviews Genetics*, 11(10):733–739, 2010.
- Dominic Masters, Josef Dean, Kerstin Klaser, Zhiyi Li, Sam Maddrell-Mander, Adam Sanders, Hatem Helal, Deniz Beker, Andrew Fitzgibbon, Shenyang Huang, et al. Gps++: Reviving the art of message passing for molecular property prediction. *arXiv preprint arXiv:2302.02947*, 2023.
- David Mendez, Anna Gaulton, A Patrícia Bento, Jon Chambers, Marleen De Veij, Eloy Félix, María Paula Magariños, Juan F Mosquera, Prudence Mutowo, Michał Nowotka, et al. ChEMBL: towards direct deposition of bioassay data. *Nucleic acids research*, 47(D1): D930–D940, 2019.
- Oscar Méndez-Lucio, Christos Nicolaou, and Berton Earnshaw. Mole: a molecular foundation model for drug discovery, 2022.
- Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.
- Hilary S Parker and Jeffrey T Leek. The practical effect of batch on genomic prediction. *Statistical applications in genetics and molecular biology*, 11(3), 2012.
- Ben Poole, Sherjil Ozair, Aaron Van Den Oord, Alex Alemi, and George Tucker. On variational bounds of mutual information. In *International Conference on Machine Learning*, pp. 5171–5180. PMLR, 2019.
- Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. 2019.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pp. 8748–8763. PMLR, 2021.



- Hubert Ramsauer, Bernhard Schöfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Thomas Adler, Lukas Gruber, Markus Holzleitner, Milena Pavlović, Geir Kjetil Sandve, et al. Hopfield networks is all you need. *arXiv preprint arXiv:2008.02217*, 2020.
- David Rogers and Mathew Hahn. Extended-connectivity fingerprints. *Journal of chemical information and modeling*, 50(5):742–754, 2010.
- Yu Rong, Yatao Bian, Tingyang Xu, Weiyang Xie, Ying Wei, Wenbing Huang, and Junzhou Huang. Self-supervised graph transformer on large-scale molecular data, 2020.
- Ana Sanchez-Fernandez, Elisabeth Rumetshofer, Sepp Hochreiter, and Günter Klambauer. Cloome: contrastive learning unlocks bioimaging databases for queries with chemical structures. *Nature*, 2023.
- Christoph Schuhmann, Romain Beaumont, Richard Vencu, Cade Gordon, Ross Wightman, Mehdi Cherti, Theo Coombes, Aarush Katta, Clayton Mullis, Mitchell Wortsman, et al. Laion-5b: An open large-scale dataset for training next generation image-text models. *Advances in Neural Information Processing Systems*, 35:25278–25294, 2022.
- Jaak Simm, Günter Klambauer, Adam Arany, Marvin Steijaert, Jörg Kurt Wegner, Emmanuel Gustin, Vladimir Chupakhin, Yolanda T Chong, Jorge Vialard, Peter Buijnsters, et al. Repurposing high-throughput image assays enables biological activity prediction for drug discovery. *Cell chemical biology*, 25(5):611–618, 2018.
- Kihyuk Sohn. Improved deep metric learning with multi-class n-pair loss objective. In *Proceedings of the 30th International Conference on Neural Information Processing Systems, NIPS’16*, pp. 1857–1865, Red Hook, NY, USA, 2016. Curran Associates Inc. ISBN 9781510838819.
- Charlotte Soneson, Sarah Gerster, and Mauro Delorenzi. Batch effect confounding leads to strong bias in performance estimates obtained by cross-validation. *PloS one*, 9(6):e100335, 2014.
- Rakshith Sharma Srinivasa, Jaejin Cho, Chouchang Yang, Yashas Malur Saidutta, Ching-Hua Lee, Yilin Shen, and Hongxia Jin. Cwcl: Cross-modal transfer with continuously weighted contrastive loss. *Advances in Neural Information Processing Systems*, 36, 2023.
- Maciej Sypetkowski, Morteza Rezanejad, Saber Saberian, Oren Kraus, John Urbanik, James Taylor, Ben Mabey, Mason Victors, Jason Yosinski, Alborz Rezazadeh Sereshkeh, et al. Rrx1: A dataset for evaluating experimental batch correction methods. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4284–4293, 2023.
- Maciej Sypetkowski, Frederik Wenkel, , Farimah Poursafaei, Nia Dickson, Karush Suri, Philip Fradkin, and Dominique Beaini. On the scalability of foundational models for molecular graphs. *arxiv*, 2024.
- Fabien Vincent, Arsenio Nueda, Jonathan Lee, Monica Schenone, Marco Prunotto, and Mark Mercola. Phenotypic drug discovery: recent successes, lessons learned and new directions. *Nature Reviews Drug Discovery*, 21(12):899–914, 2022.
- Richard Maurice Walmsley and Nicholas Billinton. How accurate is in vitro prediction of carcinogenicity? *British Journal of Pharmacology*, 162(6):1250–1258, February 2011. ISSN 1476-5381. doi: 10.1111/j.1476-5381.2010.01131.x. URL <http://dx.doi.org/10.1111/j.1476-5381.2010.01131.x>.
- Ronald Xie, Kuan Pang, Gary D. Bader, and Bo Wang. Maester: Masked autoencoder guided segmentation at pixel resolution for accurate, self-supervised subcellular structure recognition. In *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, June 2023. doi: 10.1109/cvpr52729.2023.00321. URL <http://dx.doi.org/10.1109/CVPR52729.2023.00321>.
- Sheheryar Zaidi, Michael Schaarschmidt, James Martens, Hyunjik Kim, Yee Whye Teh, Alvaro Sanchez-Gonzalez, Peter Battaglia, Razvan Pascanu, and Jonathan Godwin. Pre-training via denoising for molecular property prediction, 2022.

Xiaohua Zhai, Xiao Wang, Basil Mustafa, Andreas Steiner, Daniel Keysers, Alexander Kolesnikov, and Lucas Beyer. Lit: Zero-shot transfer with locked-image text tuning, 2022.

Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 11975–11986, 2023.

Yuanyi Zhong, Haoran Tang, Junkun Chen, Jian Peng, and Yu-Xiong Wang. Is self-supervised learning more robust than supervised learning? *arXiv preprint arXiv:2206.05259*, 2022.

Gengmo Zhou, Zhifeng Gao, Qiankun Ding, Hang Zheng, Hongteng Xu, Zhewei Wei, Linfeng Zhang, and Guolin Ke. Uni-mol: A universal 3d molecular representation learning framework. In *The Eleventh International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=6K2RM6wVqKu>.

## A ASSUMPTION OF THE INITIAL CELL STATE

There is an important distinction between phenomics - molecule and text - image contrastive training although there are initial similarities. In the text - image domain the two modalities are directly generated by the same latent variable which is the underlying semantic class. Whereas in phenomics - molecule, the observed phenomics variable is actually conditioned on molecular structure and the initial state. There are two important conclusions from this: (1) This indicates that if molecular structure has no effect on the initial cell state, there will not be a positive pairing between the molecular structure and morphological patterns captured by phenomics, making it indistinguishable from a control image. (2) There is an underlying assumption that the initial cell state  $x_i^0$  is constant. In accordance with this assumption we utilize experiments with a fixed cell line, *HUVEC-19*, and a constant genetic background. Future works can relax this assumption by taking into account phenomics experiments of the cells prior to the perturbation. This can allow the models to generalize beyond a single cell line and to diverse genetic backgrounds.

## B DATASET

Models have been trained using our in house training set and we have conducted our evaluation on two novel datasets and an open-source molecule dataset [Fay et al. \(2023\)](#):

- **Training Set:** Our training dataset comprises 1,316,283 pairs of molecules and concentration combinations, complemented by fluorescent microscopy images generated through over 2,150,000 phenomic experiments.
- **Evaluation set 1 - Unseen Images + Seen Molecules:** The first set consists of unseen images and seen molecules. Unseen microscopy images are associated with 15,058 pairs of molecules and concentrations from the training set and selected randomly.
- **Evaluation set 2 - Unseen Images + Unseen Molecules:** The second set includes previously unseen molecules, and images (consisting of 45,771 molecule and concentration pairs). Predicting molecular identities of previously unseen molecular perturbations corresponds to zero-shot prediction. Scaffold splitting was used to split this validation dataset from training ensuring minimal information leakage.
- **Evaluation set 3 - Unseen Dataset:** Finally, we utilize the RXRX3 dataset [Fay et al. \(2023\)](#), an open-source out of distribution (OOD) dataset consisting of 6,549 novel molecule and concentration pairs associated with phenomic experiments. The distribution of molecular structures differs from previous datasets, making this a challenging zero-shot prediction task.

### B.1 CONCENTRATION DETAILS

Additional details regarding the number of molecules at significant concentrations of each evaluation set are available in Table 6.

Table 6: Separated number of molecules for different concentrations at various pvalue cut-offs

Concentration	pvalue=1.0			pvalue=.1			pvalue=.01		
	Unseen Im.	Unseen Im. + Mol.	Unseen Data	Unseen Im.	Unseen Im. + Mol.	Unseen Data	Unseen Im.	Unseen Im. + Mol.	Unseen Data
.1	1497	1109	0	387	170	0	161	68	0
.25	1775	1111	1638	600	203	237	334	121	165
1.0	2721	11392	1639	1259	734	390	672	390	268
2.5	1787	4018	1636	1329	644	516	929	413	375
3.0	74	10454	0	12	1540	0	4	729	0
5.0	3	50	0	0	27	0	0	20	0
1.0	2712	11392	1636	2544	8117	792	2116	4815	625
25.0	0	2916	0	0	1734	0	0	950	0
Unique molecules	3026	14256	1639	2729	9857	823	2309	5778	642

## C IMPLEMENTATION DETAILS

In our experiments we report the top 1% recall metric as it is agnostic to the size of the dataset used. Across different datasets, top 1 metric can correspond to varying levels of difficulty due to the number of negatives evaluated. Top 1% can be used to compare models with different batch sizes, datasets, and evaluations with different number of concentrations.

### C.1 HYPERPARAMETERS

Our design choices and utilized hyperparameters for is presented in Table 7. We set batch size to 512 through experiments presented in top section of Table 1 and Figure 4 since training CLOOME model on images is not efficient compared to using pretrained models. In addition, results presented at bottom section of Table 1 are based on the best parameters found through described ablation studies (section D.4).

Table 7: Hyperparameter values utilized in our proposed MolPhenix training framework.

Parameter	Value
number of seeds	3
learning rate	1e-3
weight decay	3e-3
optimizer	AdamW
training batch size	8192
validation batch size	12000
embedding dim	512
model size	medium (38.7 M)
model structure	6 ResNet Blocks + 1 Linear layer + 1 ResNet Block + 1 Linear layer
epochs	100
self similarity clip val	.75
learnable temperature initialization	2.302
learnable bias initialization	-1.0
Distance function	arctangent of l2 distance

### C.2 RESOURCE COMPUTATION

We utilized an NVIDIA A100 GPU to train Molphenix using Ph-1 and Mol-1 embeddings, which takes approximately  $\sim 4.75$  hours each. For loss comparison experiments, we run each model using 3 different seeds and 8 different losses, resulting in a total of 114 hours of GPU processing time. For concentration experiments we train 7 runs, one for each concentration, with 3 seeds each totaling 21 runs per set of parameters. With 25 sets of parameters evaluated (11), that amounts to 2,500 A100 compute hours. Moreover, we employed 8 NVIDIA A100 GPUs to train CLOOME model on phenomics images, with an average of 40 hour usage per run. Across three seeds, that amounts to  $\sim 1000$  hours of A100 GPU usage (8 GPUs for 40 hours 3 times).

Note that, without accounting for the time to train Ph-1, MolPhenix is  $8.4 \times$  faster than the CLOOME baseline.

### C.3 S2L DISTANCE FUNCTION

To calculate inter sample distances, we utilize arctangent of l2 distances between Ph-1 embeddings. More specifically, we calculate distances with

$$\arctan(\|z_{\mathbf{x}_i} - z_{\mathbf{x}_j}\|_2^2 / c) * \frac{4}{\pi} - 1, \quad (5)$$

where  $c$  is a constant indicating the median l2 distance between a null set of embeddings. Empirically, we’ve found that setting similarities below a threshold  $k$  to 0 improves model performance:  $\lceil w \rceil^k$ .

Usage of arctan-l2 distances is motivated by an observation that cosine similarities do not effectively separate inactive molecules from other molecular pairs (Figure 6). To alleviate inactive molecule challenge, we require significant separation of CDF curves of inactive perturbations (p value  $> .9$ ) and active molecules (p  $< .01$ ). We observe that in both the plots

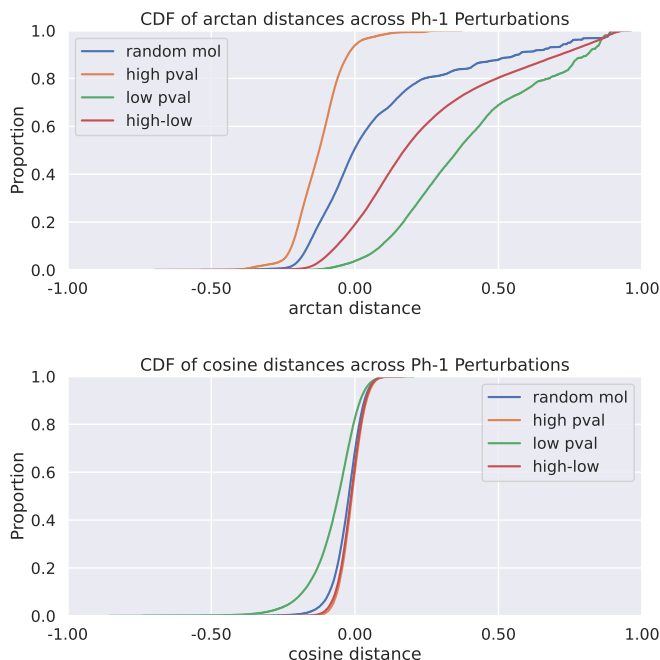


Figure 6: Plotted are cumulative densities of distance metrics for cosine similarity and arctangent of l2 distances between embeddings. Random mol corresponds to Ph-1 distances between random molecules, high pval corresponds to distances between molecules with high p-values, low pval corresponds to distances between active molecules with low p-values, finally high-low corresponds to distances between active and inactive molecules.

using arctangent and cosine similarities achieves this purpose. However, if we compare high p-value curves with high-low, we find that in the case of cosine similarities they are almost identical. This indicates that the distribution of cosine similarities between active - inactive molecules is almost identical to that of inactive - inactive molecules. In contrast, when using arctangent similarities, we observe that the two CDF curves are well separated.

This property of l2 distances can inform our model training to identify inactive-inactive molecules. These results informed our decision to utilize arctangent of l2 distances to specify sample similarities for the S2L loss.

## D ADDITIONAL RESULTS

### D.1 PREDICTING MOLECULAR ACTIVITY

Given the significance of identifying active molecules, we evaluate the ability of the chemical encoder to predict molecular activity. To do so, we assessed whether embeddings generated from the chemical encoder can be used to predict calculated p-values for unseen molecules. We fit a logistic regression on molecular embeddings from the training set, classifying whether a molecular perturbation and concentration have a p-value below .01. We find that the trained logistic regression is capable of predicting molecular activity on two downstream datasets with a non-overlapping set of molecules, Figure 8. In addition, we provide a u-map of molecular embedding for the unseen dataset RXX3, colored by p-value. We qualitatively observe a clustering of active molecules using a U-map (Figure 7). It demonstrates that predicting compounds activity is possible using MolPhenix chemical encoder as molecules representations are distinct, independent of the experimented dosage concentration.

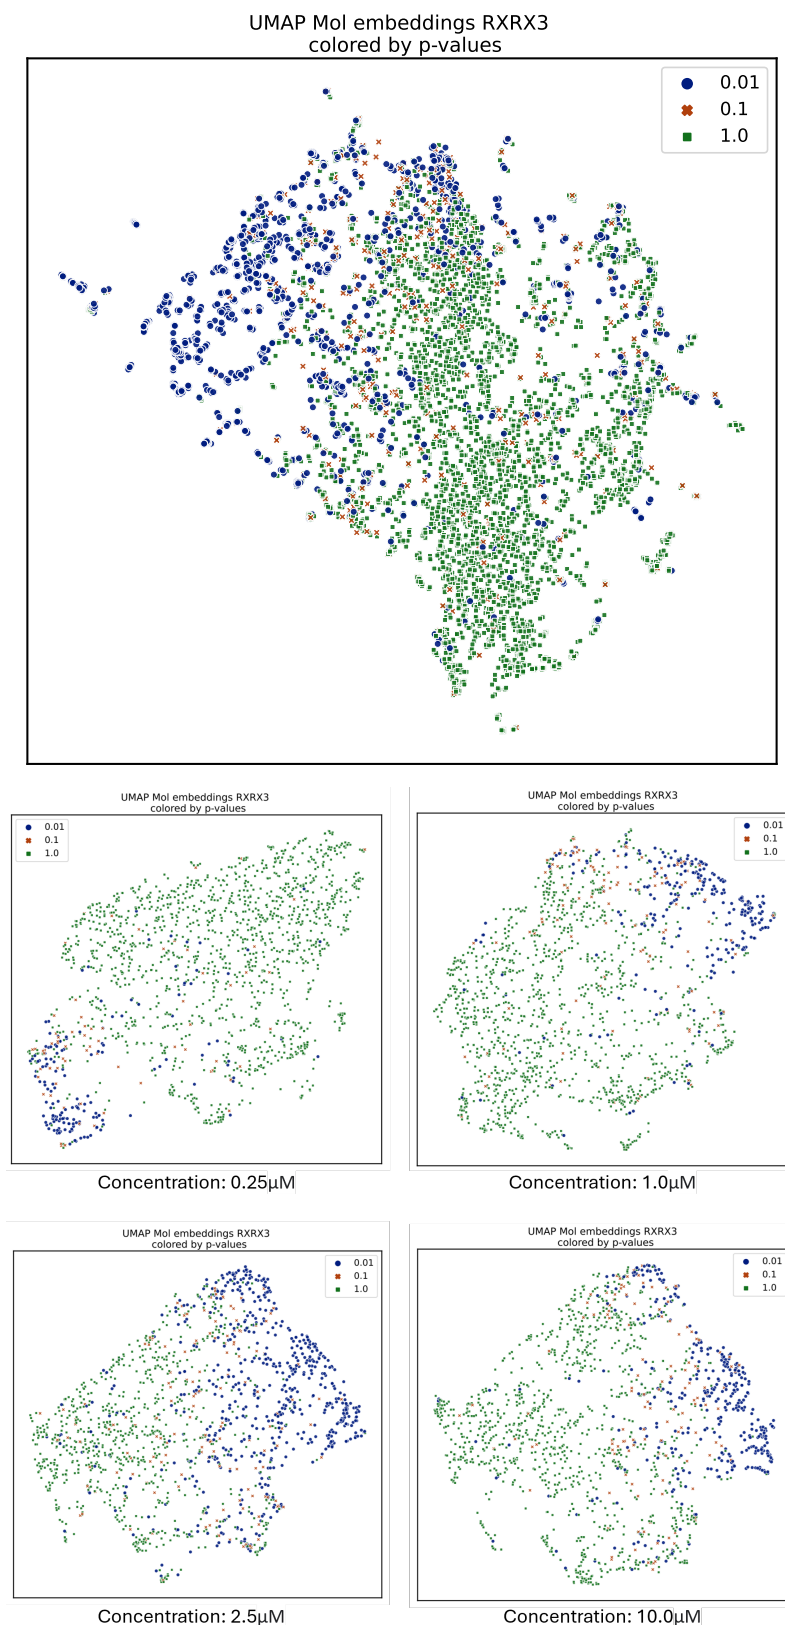


Figure 7: U-map demonstrating dimensionality reduction of the chemical embeddings of unseen dataset RXRX3. First two dimensions are visualized and points are colored corresponding to their activity measured in phenomics experiments. Activity is evaluated using p-values calculated using technical replicability of Ph-1 embeddings. Top plot shows the u-map figure of all chemical embeddings, and bottom figure contains u-map figure of representations at specific concentrations.






Figure 8: **Top left:** ROC AUC of logistic regression predicting molecular activity on new dataset. **Top right:** ROC AUC of logistic regression predicting molecular activity on validation dataset with new molecules and new images.

## D.2 ZERO SHOT BIOLOGICAL VALIDATION


We conduct a preliminary investigation into whether MolPhenix can be used to identify biological relationships without the need for conducting the underlying experiments. To this end, we evaluate on a subset of ChEMBL with curated pairs of gene knockouts and molecular perturbants [Mendez et al. \(2019\)](#). These pairs of perturbations were curated due to the similarity of their effects on cells, although these might not always be captured through phenomic experiments. Thus, there is maximum performance that can be reached through just phenomic data, which we assume to be achieved by experimental data embedded using Ph-1.

To evaluate MolPhenix’s ability to identify previously known biological associations directly from data, we embed phenomics experiments from gene knockouts using the vision encoder. To perform in-silico screening, we then embed the molecular structures associated with positive pairs using the chemical encoder. Generating molecular embeddings and the corresponding concentrations does not utilize any experimental data. We then calculate cosine similarities between embeddings of phenomics experiments evaluating gene knockouts, and representations of the chemical representations along with encoded concentrations. Using the computed cosine similarities we are then able to assess whether MolPhenix is capable of identifying known associations between gene knockouts and molecular structures. Since there is no information on molecular concentration at which the cells must be treated with, we repeat the experiment across 4 concentrations. To get a null distribution of cosine similarities we take pairs of genes knockouts and molecules for which there are no annotated relationships. We calculate a cut-off for a low and high percentiles, and then evaluate what percentage of pairs of genes and molecules with known relationships exceed the set thresholds.

Figure 9 demonstrates that in-silico screening using MolPhenix Molecular encoder is capable of recovering a significant portion of known interactions. This is performed without the use of experimental data on the molecular encoder. It is difficult to estimate an upper bound on the expected performance due to uncertainty in the quality of curation of known pairs, presence of unknown associations between genes and molecules, and uncertainty regarding molecular concentration. There is a clear trend however that MolPhenix molecular encoder is capable of recovering a meaningful fraction of these interactions.

**Table 8: Evaluation on cumulative concentrations for active molecules:** Average Top-1% and Top-5% recall accuracies of methods utilizing different contrastive learning loss functions and concentration encoding information. We evaluate all methods on unseen images, unseen images and unseen molecules and an unseen dataset for zero-shot retrieval. Entries in **bold** denote best performance when the loss function is fixed while entries in **highlight** denote best performance across all guidelines.

Method	Explicit Concentration (ours)	Modality	top-1%				top-5%			
			Unseen Images	Unseen Images + Unseen Molecules	Unseen Dataset (0-shot)	Avg.	Unseen Images	Unseen Images + Unseen Molecules	Unseen Dataset (0-shot)	Avg.
CLIP	<b>X</b>	Ph-1	.3373	<b>.4228</b>	.3514	.3038	<b>.6162</b>	<b>.7182</b>	<b>.3660</b>	.5668
Hopfield-CLIP	<b>X</b>	Ph-1	.2578	.3559	.1256	.2464	.5457	<b>.6751</b>	.3270	.5159
InfoLOOB	<b>X</b>	Ph-1	.3351	<b>.4206</b>	.1563	.3040	<b>.6128</b>	<b>.7204</b>	<b>.3730</b>	.5687
CLOOME	<b>X</b>	Ph-1	.3572	.4348	.1658	.3193	.6330	.7259	.3918	.5836
CLOOME	sigmoid	Ph-1	.5813	.4968	.2360	.4380	.8748	.7658	.4859	.7088
CLOOME	logarithm	Ph-1	<b>.6057</b>	<b>.5255</b>	<b>.2445</b>	<b>.4586</b>	<b>.8858</b>	.8117	<b>.4957</b>	<b>.7310</b>
CLOOME	one-hot	Ph-1	.5967	.5255	.2380	.4534	.8800	<b>.8120</b>	.4829	.7250
DCL	<b>X</b>	Ph-1	.6363	.6177	.3184	.5241	.8638	.8180	.5632	.7483
DCL	sigmoid	Ph-1	.8858	.6694	.4527	.6693	<b>.9600</b>	.8472	.6845	.8305
DCL	logarithm	Ph-1	.8934	.6952	.4511	.6799	<b>.8788</b>	<b>.6889</b>	<b>.8419</b>	<b>.8419</b>
DCL	one-hot	Ph-1	<b>.8901</b>	<b>.7002</b>	<b>.4601</b>	<b>.6834</b>	.8770	.6873	.8411	
CWCL	<b>X</b>	Ph-1	.7091	.6529	.3556	.5725	.9018	.8368	.6027	.7804
CWCL	sigmoid	Ph-1	.9138	.6985	.4810	.6977	.9681	.8643	.7070	.8464
CWCL	logarithm	Ph-1	.9141	.7248	.4815	.7068	.9651	.8920	<b>.7131</b>	<b>.8567</b>
CWCL	one-hot	Ph-1	<b>.9128</b>	<b>.7261</b>	<b>.4850</b>	<b>.7079</b>	<b>.9665</b>	<b>.8927</b>	.6998	.8530
SigLip	<b>X</b>	Ph-1	.7763	.6401	.3396	.5853	.9361	.83038	.5714	.7792
SigLip	sigmoid	Ph-1	.9463	.6931	.4576	.6990	.9816	.8606	.6759	.8393
SigLip	logarithm	Ph-1	<b>.9493</b>	.7256	.4868	.7205	.9814	<b>.8926</b>	.7019	<b>.8586</b>
SigLip	one-hot	Ph-1	.9489	.7210	<b>.4859</b>	.7186	<b>.9823</b>	.8868	.7045	.8578
MolPhenix (ours)	<b>X</b>	Ph-1	.9097	.6759	.4181	.6679	.9768	.8539	.6436	.8247
MolPhenix (ours)	sigmoid	Ph-1	.9423	.7155	.4573	.7050	.9808	.8775	.6778	.8453
MolPhenix (ours)	logarithm	Ph-1	.9426	.7451	.4727	.7201	.9808	.8964	.6952	.8574
MolPhenix (ours)	one-hot	Ph-1	<b>.9430</b>	<b>.7490</b>	<b>.4850</b>	<b>.7256</b>	<b>.9816</b>	<b>.8984</b>	.7040	<b>.8613</b>
MolPhenix (ours)	<b>X</b>	Ph-1 + Mol-1	.9105	.6710	.4501	.6772	.9755	.8527	.7098	.8460
MolPhenix (ours)	sigmoid	Ph-1 + Mol-1	.9395	.7034	.5252	.7227	.9811	.8729	.7630	.8723
MolPhenix (ours)	logarithm	Ph-1 + Mol-1	.9413	.7505	.5473	.7463	.9811	<b>.9085</b>	<b>.7878</b>	<b>.8924</b>
MolPhenix (ours)	one-hot	Ph-1 + Mol-1	<b>.9430</b>	<b>.7514</b>	<b>.5577</b>	<b>.7507</b>	<b>.9830</b>	.9043	.7821	.8898



**Figure 9: Evaluation of 0-shot ChEMBL identification of gene knockout and molecular phenomic similarities.** On the X axis are percentile ranges, at which points the threshold is computed for cosine similarities. On the y axis is plotted total recall of recovered known interactions. Grey x plotted for each range indicate baseline recall. Orange line indicates MolPhenix-Molecular encoding of chemical compounds and MolPhenix-Vision for encoding gene knockout phenomics experiment. Blue line indicates Ph-1 encoding of phenomics experiments for both the molecular perturbation and gene knockouts. In-silico encoding of molecular perturbation, as well as the corresponding concentration, recovers a significant fraction of observed interactions.

### D.3 ADDRESSING CHALLENGES IN CONTRASTIVE PHENOMIC RETRIEVAL

Table 8 and 10 show the complete Top 1% and 5% results of evaluation on cumulative concentrations on only active and all molecules, respectively. Similarly, Table 9 and 11 presents the full retrieval results of held-out concentrations experiments. In comparison to prior loss functions, S2L loss objective demonstrates consistent high retrieval rate in all tasks and molecular groups (i.e. all or active molecules), while using the same modality (Ph-1) and with or without explicit concentration information.

### D.4 ABLATION STUDIES

Figure 10 and Table 13, 14, 15, 16 and 17 present top-1% recall accuracy across for the full ablation study on the variation of MolPhenix key components. We note that compact

Table 9: **Evaluation on held-out concentration for active molecules:** Average Top-1% and Top-5% recall accuracies of methods utilizing different contrastive learning loss functions and concentration encoding information. We evaluate all methods on unseen images, unseen images and unseen molecules and an unseen dataset for zero-shot retrieval. Entries in **bold** denote highest performance when the loss function is fixed while entries in **highlight** denote highest performance across all guidelines.

Method	Explicit Concentration (ours)	Modality	top-1%				top-5%			
			Unseen Images	Unseen Images + Unseen Molecules	Unseen Dataset (o-shot)	Avg.	Unseen Images	Unseen Images + Unseen Molecules	Unseen Dataset (o-shot)	Avg.
CLIP	<b>X</b>	Ph-1	<b>.2109</b>	<b>.2425</b>	<b>.1519</b>	<b>.2018</b>	<b>.4458</b>	<b>.4968</b>	<b>.3591</b>	<b>.4339</b>
Hopfield-CLIP	<b>X</b>	Ph-1	<b>.1581</b>	<b>.2034</b>	<b>.1198</b>	<b>.1604</b>	<b>.3783</b>	<b>.4413</b>	<b>.3045</b>	<b>.3747</b>
InfoLOOB	<b>X</b>	Ph-1	<b>.2122</b>	<b>.2496</b>	<b>.1501</b>	<b>.2040</b>	<b>.4443</b>	<b>.5003</b>	<b>.3515</b>	<b>.4320</b>
CLOOME	<b>X</b>	Ph-1	.2164	.2461	.1479	.2035	.4590	.4956	.3528	.4358
CLOOME	sigmoid	Ph-1	.3338	.2681	.1801	.2606	.6037	.5202	.3879	.5039
CLOOME	logarithm	Ph-1	.3094	.2345	.1665	.2368	.5960	.4874	.3534	.4790
CLOOME	one-hot	Ph-1	.3073	.2040	.1670	.2261	.5997	.4246	.3657	.4633
DCL	<b>X</b>	Ph-1	.4717	.4027	.2841	.3861	.7352	.6579	.5322	.6417
DCL	sigmoid	Ph-1	<b>.7282</b>	<b>.4100</b>	<b>.3560</b>	<b>.4980</b>	<b>.9226</b>	<b>.6561</b>	<b>.6015</b>	<b>.7267</b>
DCL	logarithm	Ph-1	.6903	.3558	.3211	.4557	.8869	.6146	.5667	.6894
DCL	one-hot	Ph-1	.6562	.3607	.3272	.4480	.8831	.5983	.5659	.6824
CWCL	<b>X</b>	Ph-1	.5731	.4403	.3232	.4455	.8218	.6833	.5706	.6919
CWCL	sigmoid	Ph-1	.7780	<b>.4425</b>	<b>.3777</b>	<b>.5327</b>	<b>.9386</b>	<b>.6844</b>	<b>.6244</b>	<b>.7491</b>
CWCL	logarithm	Ph-1	.7452	.3989	.3523	.4988	.9117	.6482	.5962	.7187
CWCL	one-hot	Ph-1	.7048	.4009	.3593	.4883	.9037	.6284	.6061	.7127
SigLip	<b>X</b>	Ph-1	.5718	.4217	.3021	.4318	.8104	.6602	.5176	.6627
SigLip	sigmoid	Ph-1	<b>.8166</b>	<b>.4640</b>	<b>.3830</b>	<b>.5612</b>	<b>.9623</b>	<b>.7023</b>	<b>.6080</b>	<b>.7575</b>
SigLip	logarithm	Ph-1	.8097	.4391	.3747	.5411	.9437	.6746	.6046	.7409
SigLip	one-hot	Ph-1	.7561	.4020	.3345	.4975	.9279	.6248	.5557	.7028
MolPhenix (ours)	<b>X</b>	Ph-1	<b>.8334</b>	.4615	<b>.3792</b>	.5580	<b>.9638</b>	.6937	<b>.6128</b>	.7567
MolPhenix (ours)	sigmoid	Ph-1	.8356	<b>.4692</b>	.3765	<b>.5571</b>	<b>.9638</b>	<b>.7068</b>	.6115	<b>.7607</b>
MolPhenix (ours)	logarithm	Ph-1	.7953	.4466	.3664	.5361	.9466	.6889	.5924	.7426
MolPhenix (ours)	one-hot	Ph-1	.7489	.4088	.3379	.4985	.9325	.6465	.5644	.7144
MolPhenix (ours)	<b>X</b>	Ph-1 & Mol-1	<b>.8277</b>	.4739	.4071	.5695	<b>.9602</b>	.7041	.6798	.7813
MolPhenix (ours)	sigmoid	Ph-1 & Mol-1	.8218	<b>.4771</b>	.4287	<b>.5758</b>	.9588	.7117	<b>.7045</b>	<b>.7916</b>
MolPhenix (ours)	logarithm	Ph-1 & Mol-1	.7836	.4757	<b>.4297</b>	.563	.9402	<b>.7138</b>	.7011	.7850
MolPhenix (ours)	one-hot	Ph-1 & Mol-1	.7391	.4307	.3940	.5212	.9198	.6724	.6698	.7540

Table 10: **Evaluation on cumulative concentrations for active and inactive perturbations** Average Top-1% and Top-5% Recall accuracy of methods utilizing different contrastive learning methods. Best performing methods are highlighted in **bold**.

Loss	Explicit Concentration	Modality	top-1%				top-5%			
			Unseen Images	Unseen Images + Unseen Molecules	Unseen Dataset (o-shot)	Avg.	Unseen Images	Unseen Images + Unseen Molecules	Unseen Dataset (o-shot)	Avg.
CLIP	<b>X</b>	Ph-1	.1761	.1867	.0734	.1454	.3710	.3769	.2065	.3181
Hopfield-CLIP	<b>X</b>	Ph-1	.1531	.1709	.0673	.1304	.3464	.3637	.1942	.3014
InfoLOOB	<b>X</b>	Ph-1	.1746	.1860	.0745	.1450	.3697	.3756	.2058	.3170
CLOOME	<b>X</b>	Ph-1	.1968	.2005	.0911	.1628	.3938	.3888	.2321	.3383
CLOOME	sigmoid	Ph-1	.3875	.2592	.1415	.2627	.5662	.4601	.2940	.4401
CLOOME	logarithm	Ph-1	<b>.4088</b>	.3046	<b>.1503</b>	.2879	.5730	.5166	.3053	.4650
CLOOME	one-hot	Ph-1	.4080	<b>.3123</b>	.1496	.2900	<b>.5801</b>	<b>.5306</b>	<b>.3054</b>	<b>.4720</b>
DCL	<b>X</b>	Ph-1	.3277	.2562	.1364	.2401	.4856	.4170	.2768	.3931
DCL	sigmoid	Ph-1	.4881	.3380	.2009	.3423	.6222	.5186	.3381	.4930
DCL	logarithm	Ph-1	.4983	.3615	.2122	.3573	.6311	.5581	.3587	.5160
DCL	one-hot	Ph-1	<b>.5226</b>	<b>.3790</b>	<b>.2288</b>	<b>.3768</b>	<b>.6791</b>	<b>.5870</b>	<b>.3968</b>	<b>.5543</b>
CWCL	<b>X</b>	Ph-1	.3635	.2696	.1526	.2619	.5122	.4267	.2933	.4107
CWCL	sigmoid	Ph-1	.5070	.3457	.2101	.3542	.6378	.5272	.3462	.5037
CWCL	logarithm	Ph-1	.5146	<b>.3725</b>	.2246	.3706	.6437	.5733	.3660	.5277
CWCL	one-hot	Ph-1	<b>.5401</b>	<b>.3849</b>	<b>.2336</b>	<b>.3862</b>	<b>.6882</b>	<b>.5991</b>	<b>.4001</b>	<b>.5625</b>
SigLip	<b>X</b>	Ph-1	.3729	.2544	.1470	.2581	.5200	.4179	.2838	.4072
SigLip	sigmoid	Ph-1	.5021	.3275	.2072	.3456	.6360	.5231	.3430	.5007
SigLip	logarithm	Ph-1	.5156	.3636	.2233	.3675	.6452	.5689	.3653	.5265
SigLip	one-hot	Ph-1	<b>.5354</b>	<b>.3745</b>	<b>.2317</b>	<b>.3805</b>	<b>.6858</b>	<b>.5928</b>	<b>.3945</b>	<b>.5577</b>
S2L (ours)	<b>X</b>	Ph-1	.4688	.2852	.1838	.3126	.5970	.4519	.3171	.4554
S2L (ours)	sigmoid	Ph-1	.5071	.3441	.2144	.3552	.6428	.5315	.3554	.5099
S2L (ours)	logarithm	Ph-1	.5183	.3700	.2275	.3720	.6492	.5650	.3756	.5300
S2L (ours)	one-hot	Ph-1	<b>.5433</b>	<b>.3819</b>	<b>.2384</b>	<b>.3879</b>	<b>.6954</b>	<b>.5895</b>	<b>.4030</b>	<b>.5626</b>
S2L (ours)	<b>X</b>	Ph-1	.4688	.2729	.2001	.3139	.5956	.4374	.3430	.4587
S2L (ours)	sigmoid	Ph-1 & Mol-1	.4983	.3230	.2397	.3537	.6343	.5035	.3790	.5056
S2L (ours)	logarithm	Ph-1 & Mol-1	.5101	.3589	.2535	.3742	.6398	.5660	.3992	.5350
S2L (ours)	one-hot	Ph-1 & Mol-1	.5370	.3720	<b>.2676</b>	<b>.3922</b>	<b>.6870</b>	<b>.5888</b>	<b>.4326</b>	<b>.5695</b>

**Table 11: Evaluation on held-out concentrations for active and inactive perturbations**  
 Average Top-1% and Top-5% Recall accuracy of methods utilizing different contrastive learning methods. Best performing methods are highlighted in **bold**.

Loss	Explicit Concentration	Modality	top-1%				top-5%			
			Unseen Images	Unseen Images + Unseen Molecules	Unseen Dataset (o-shot)	Avg.	Unseen Images	Unseen Images + Unseen Molecules	Unseen Dataset (o-shot)	Avg.
CLIP	<b>X</b>	Ph-1	.1684	.1111	.0964	.1253	.3916	.2545	.2356	.2476
Hopfield-CLIP	<b>X</b>	Ph-1	.1290	.0921	.0756	.0989	.3485	.2287	.2095	.2217
InfoLOOB	<b>X</b>	Ph-1	.1715	.1114	.0948	.1259	.3944	.2578	.2349	.2495
CLOOME	<b>X</b>	Ph-1	.1745	.1088	.0910	.1248	.4093	.2487	.2355	.2439
CLOOME	sigmoid	Ph-1	<b>.2573</b>	<b>.1208</b>	<b>.1062</b>	<b>.1614</b>	<b>.5169</b>	<b>.2638</b>	<b>.2513</b>	<b>.3440</b>
CLOOME	logarithm	Ph-1	.2379	.1081	.0992	.1484	.4958	.2444	.2324	.3242
CLOOME	one-hot	Ph-1	.2346	.0970	.0974	.1430	.5014	.2224	.2348	.3195
DCL	<b>X</b>	Ph-1	.3516	.1655	.1533	.2235	.5693	.3125	.3006	.3082
DCL	sigmoid	Ph-1	<b>.4741</b>	<b>.1725</b>	<b>.1726</b>	<b>.2731</b>	<b>.6637</b>	<b>.3261</b>	<b>.3105</b>	<b>.3204</b>
DCL	logarithm	Ph-1	.4286	.1596	.1581	.2488	.6244	.3071	.3032	.3056
DCL	one-hot	Ph-1	.4308	.1495	.1600	.2468	.6244	.2938	.3015	.2966
CWCL	<b>X</b>	Ph-1	.4126	.1801	.1667	.2531	.6128	.3266	.3066	.3194
CWCL	sigmoid	Ph-1	<b>.5112</b>	<b>.1856</b>	<b>.1811</b>	<b>.2926</b>	<b>.6901</b>	<b>.3384</b>	<b>.3190</b>	<b>.3314</b>
CWCL	logarithm	Ph-1	.4664	.1696	.1709	.2690	.6502	.3195	.3066	.3148
CWCL	one-hot	Ph-1	.4681	.1612	.1734	.2676	.6465	.3019	.3104	.3050
SigLip	<b>X</b>	Ph-1	.3942	.1578	.1390	.2303	.5931	.3015	.2737	.2914
SigLip	sigmoid	Ph-1	<b>.5392</b>	<b>.1828</b>	<b>.1710</b>	<b>.2977</b>	<b>.7102</b>	<b>.3399</b>	<b>.3121</b>	<b>.3398</b>
SigLip	logarithm	Ph-1	.5022	.1698	.1669	.2796	.6841	.3240	.3068	.3177
SigLip	one-hot	Ph-1	.4657	.1443	.1451	.2517	.6544	.2879	.2790	.2847
S2L (ours)	<b>X</b>	Ph-1	.5336	.1842	.1713	.2963	.6961	.3322	.3045	.3221
S2L (ours)	sigmoid	Ph-1	<b>.5409</b>	<b>.1899</b>	<b>.1753</b>	<b>.3020</b>	<b>.7178</b>	<b>.3469</b>	<b>.3201</b>	<b>.3372</b>
S2L (ours)	logarithm	Ph-1	.5036	.1791	.1727	.2851	.6925	.3342	.3157	.3275
S2L (ours)	one-hot	Ph-1	.4726	.1537	.1521	.2595	.6696	.2998	.2887	.2958
S2L (ours)	<b>X</b>	Ph-1 & Mol-1	.5248	.1829	.1910	.2996	.6904	.3268	.3305	.3281
S2L (ours)	sigmoid	Ph-1 & Mol-1	<b>.5338</b>	<b>.1897</b>	<b>.2029</b>	<b>.3088</b>	<b>.7098</b>	<b>.3427</b>	<b>.3495</b>	<b>.3452</b>
S2L (ours)	logarithm	Ph-1 & Mol-1	.4900	.1839	.2031	.2923	.6776	.3354	.3511	.3411
S2L (ours)	one-hot	Ph-1 & Mol-1	.4622	.1569	.1762	.2651	.6578	.3030	.3187	.3087

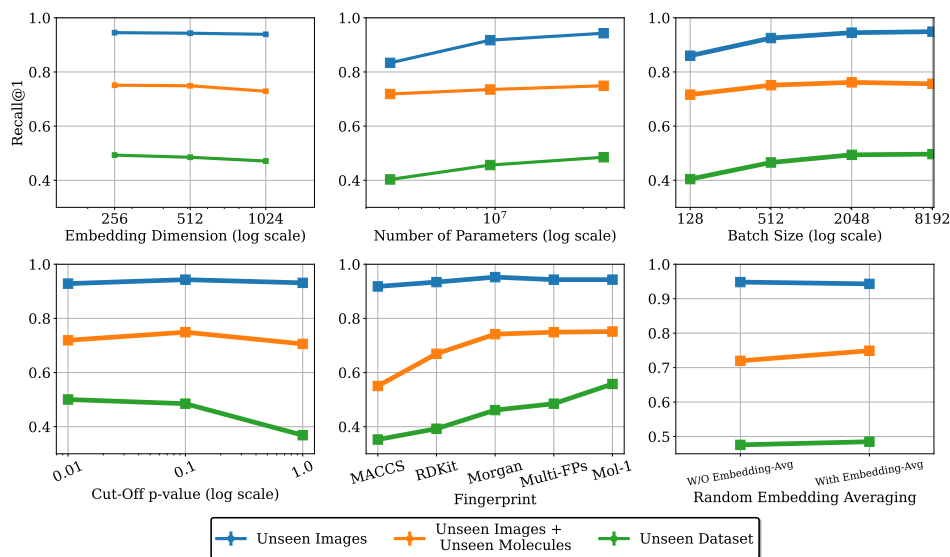


Figure 10: Ablations of top-1 % recall accuracy with (**top-left**) the size of embedding dimension, (**top-center**) number of parameters, (**top-right**) batch size, (**bottom-left**) cutoff  $p$  value, (**bottom-center**) fingerprint type, and (**bottom-right**) random batch averaging. Compact embedding sizes from pretrained models, larger number of parameters, larger batch sizes, lower cutoff  $p$ -values, pretrained Mol-1 fingerprints and presence of random batch averaging improving retrieval of our MolPhenix framework.

embedding sizes from pretrained models stabilize training. This indicates that embeddings are expressive and accurately capture intricate aspects of molecules. Larger batch sizes result in a greater number of negative samples, hence improving performance. This is in line with prior contrastive learning methods continuing to improve by increasing the batch size [Chen et al. \(2020\)](#). Increasing the number of parameters leads to more expressive models thereby enhancing retrieval performance. This result is in accordance with recent advances

Model size	Depth	Width	Unseen images	Unseen images + Unseen molecules	Unseen dataset (O-shot)
Tiny - 2.7m	4 ResBlocks	256	.8337	.7186	.4030
Small - 9.4m	6 ResBlocks	512	.9174	.7352	.4562
Medium - 38.7m	8 ResBlocks	1024	.9430	.7490	.485

Table 12: Ablations across different model sizes. Larger capacity models are found to be more expressive.

Batch size	Unseen images	Unseen images + Unseen molecules	Unseen dataset (O-shot)
128	.8600	.7163	.4044
512	.9252	.7511	.4657
2048	.9450	.7616	.4940
8192	.9489	.7563	.4966

Table 13: Ablation across different batch sizes. Larger batch sizes benefit contrastive learning.

in language modelling and scaling laws across different data and compute budgets [Kaplan et al. \(2020\)](#).

Dim size	Unseen images	Unseen images + Unseen molecules	Unseen dataset (O-shot)
256	.9452	.7510	.4929
512	.9430	.7490	.4850
1024	.9392	.7288	.4710

Table 14: Ablation across different embedding dimensions. Compact embedding sizes capture more molecular information.

cut-off	Unseen images	Unseen images + Unseen molecules	Unseen dataset (0-shot)
p < 1.0	.9312	.7057	.3686
p < .1	.9430	.7490	.4850
p < .01	.9284	.7192	.5005

Table 15: Ablation across different p-value cutoff thresholds. p values < .1 benefit retrieval of active molecules.

fingerprint	unseen images	unseen images + unseen molecule	unseen dataset
MACCS	.9180	.5503	.3526
RDKit	.9341	.6693	.3925
Morgan	.9524	.7417	.4613
Multi-FPs	.9430	.7490	.485
Ph-1 + Mol-1	.9430	.7514	.5577

Table 16: Ablation across different fingerprint types. A combination of embeddings bootstrapped from Ph-1 and Mol-1 significantly benefit retrieval.

	Unseen images	Unseen images + Unseen molecules	Unseen dataset (0-shot)
W/O Random Embedding Avg.	.9482	.7198	.4759
With Random Embedding Avg.	.9430	.7490	.485

Table 17: Ablation across random embedding averaging. Utilizing random batch averaging stabilizes training and benefits retrieval.